


PHILIP CAREY

AI and Software

 (650) 799-7709

 prcarey11@gmail.com

 [linkedin.com/in/philip-carey-596980194/](https://www.linkedin.com/in/philip-carey-596980194/)

SKILLS

Tech-Related

- **Python**
 - General: numpy, pandas, matplotlib, seaborn, altair
 - ML: Pytorch, TensorFlow, scikit-learn
- **MySQL**
- **Ruby/Ruby on Rails**
- **AWS Cloud Services**
- **Git**
- **Tableau**

Applied Statistical Methods

- **Machine Learning**
- **Multivariate Regression**
- **Classification**
- **Time Series**
- **Dimensionality Reduction**
- **Clustering**

Other/Miscellaneous

- **JIRA/Agile**
- **Technical Writing**
- **Prompt Engineering**
- **Data Annotation**
- **Data Labeling**

WORK EXPERIENCE

AI Tutor

August 2024 – Present

xAI

- Conduct large-scale evaluations of generative AI models, completing over 3000+ annotation tasks and 500+ hours of detailed data analysis
- Utilize Reinforcement Learning with Human Feedback (RLHF) to refine AI-generated responses, assessing criteria such as a model's Instruction Following, Truthfulness, Helpfulness, and Writing Style
- Develop structured evaluation criteria for chatbot responses, improving model performance through systematic prompt engineering and response grading methodologies
- Perform quality assurance on RLHF tasks, ensuring consistency across AI Tutor contributions

Machine Learning Engineer Intern *June 2023 – September 2023*

AppFolio, Inc.

Main Project (Surfacing CTA'S for inbound texts):

- Created a Python-based LLM skill to call the OpenAI API; utilized prompt engineering techniques to effectively classify incoming text message data as "response" or "no response" needed
- Built out the ETL pipeline for interpreting text message threads with a stack of Ruby on Rails backend, AWS S3 buckets/lambda functions, and Snowflake SQL queries
- Recognized design flaws in our database schema preventing the use of a simpler NLP classifier, and generated a detailed report on both what went wrong and proposed improvements for data labeling

Company-Wide Hackday Project (Pet Policy Generation):

- Collaborated with a group of 4 other team members to develop a system that utilized text from housing leases to generate pet policies for onboarding customers
- Responsible for the ML aspect of the project; created a service that called the OpenAI API, engineered a prompt to generate accurate pet policies using the lease text

Oral Communication

- Delivered two comprehensive 30-minute presentations to the engineering and product groups at the end of the internship

EDUCATION

UCLA (Los Angeles, CA)

September 2022–August 2024

Master of Science,
Applied Statistics and Data Science
Grade: 4.0/4.0

UCSB (Santa Barbara, CA)

September 2018–June 2022

Bachelor of Science,
Statistics and Data Science